

Machine Learning

Jakub Adamczyk, Faculty of Computer Science, AGH

Time series forecasting

Part 3:
neural models

Why deep learning?

Complex tasks

- primarily excel at highly complex tasks with lots of data
- **big data:**
 - high frequency and long history
- **complex tasks:**
 - complex relations, e.g. multiple variable seasonalities
 - dynamic, requiring adaptation, with changing patterns and noise
- **multivariate time series:**
 - many time series, with cross-series relations
- **long forecasting horizons:**
 - can be quite precise compared to classical methods

Direct multi-step (DMS) forecasts

- statistical models basically always perform **autoregressive forecasting**:
 - forecast 1 step ahead at a time, assume previous forecasts are true
 - this results in error accumulation and higher error bias
 - also known as iterated multi-step (IMS) forecasting (or recursive forecasting)
- neural networks are typically **multioutput**, i.e. can easily have many output neurons
- this results in **direct multi-step (DMS)** forecasts, which:
 - avoids error accumulation, but has higher error variance
 - acts as regularization, since it has to optimize many horizons at a same time

IMS vs DMS forecasting - additional resources

- [CrossValidated - Time Series One Step Ahead vs N-Step Ahead](#)
- ["Recursive and direct multi-step forecasting: the best of both worlds" S. Taieb, R. Hyndman](#)
- ["A comparison of direct and iterated multistep AR methods for forecasting macroeconomic time series" M. Marcellino et al.](#)
- ["When are Direct Multi-Step and Iterative Forecasts Identical?" T. McElroy](#)
- ["Direct Versus Iterated Multiperiod Volatility Forecasts" E. Ghysels et al.](#)
- ["An Empirical Investigation of Direct and Iterated Multistep Conditional Forecasts" M. McCracken, J. McGillicuddy](#)

Pretraining

- **transfer learning** and **pretraining** of time series recently became possible:
 - novel architectures, particularly transformers
 - available massive datasets
 - utilize previous knowledge and reduce overfitting
- **foundational models** start to emerge, with e.g. few-shot and zero-shot capabilities
- does this work?
 - many works lack fair evaluation, or fail against simple baselines
 - comparisons are often artificial and unrealistic, e.g. lots of data, highly multivariate
 - just research currently, not well tested in the industry
- definitely a **research direction** in near future

Deep learning approaches

Deep learning approaches

- **linear networks:**
 - Linear, DLinear, NLinear, RLinear etc.
 - simple, 1-layer networks for univariate time series
- **MLP-based models:**
 - N-BEATS, N-HiTS, TSMixer, TiDE etc.
 - learn complex relations and decompositions by using stacks of MLPs
- **transformers:**
 - PatchTST, Autoformer, FEDformer, Pyraformer etc.
 - pure transformer architectures, often with complex attention modifications

Deep learning approaches

- **pretrained foundational models:**
 - TimesFM, Chronos, Lag-Llama, Moirai, TimeGPT etc.
 - first really successful transfer learning for time series
- **State Space Models (SSMs):**
 - LSSL, MambaTS, Chimera, SpaceTime etc.
 - state-space models theory unifies ETS, CNNs, RNNs, and a few other things
- **graph neural networks (GNNs):**
 - T-GCN, DGSL, GaAN, STGNN etc.
 - typically used for spatio-temporal forecasting, e.g. traffic demand

Deep learning approaches

- **recurrent networks (RNNs):**
 - old (mostly obsolete), e.g. LSTM, GRU, DeepAR
 - modern, e.g. RWKV-TS, TFT, P-sLSTM
 - built for sequence prediction, fast inference, but can be hard to train
- **convolutional networks (CNNs):**
 - old (mostly obsolete), e.g. TCN, DeepTCN
 - modern, e.g. MICN, TimesNet, SCINet
 - typically based on dilated convolutions and causal convolutions
- classical ones are generally obsolete, but modern ones are noteworthy

Agenda

- we will go over **representative architectures** from the most commonly used groups:
 - linear models
 - MLP-based
 - transformers
- lastly, we will cover **important research direction** - pretrained foundation models
- we omit others, because:
 - SSMs are not well proven or popular (yet)
 - GNNs are specific for spatio-temporal forecasting
 - RNNs and CNNs are mostly obsolete (with some notable exceptions)

Linear networks

Linear

"Are Transformers Effective for Time Series Forecasting?" A. Zeng et al.

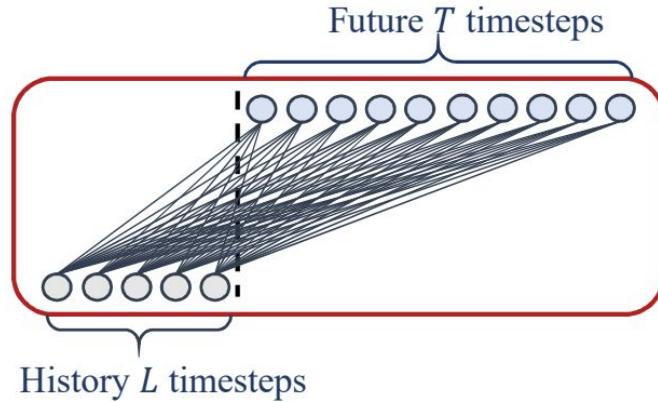
- **LTSF-Linear** (Long-term Time Series Forecasting Linear)
- just a linear projection from L to T values

$$\hat{y}_i = W y_i \quad W \in \mathbb{R}^{T \times L}$$

L - lookback, number of previous steps (hyperparam.)

T - how many steps to forecast

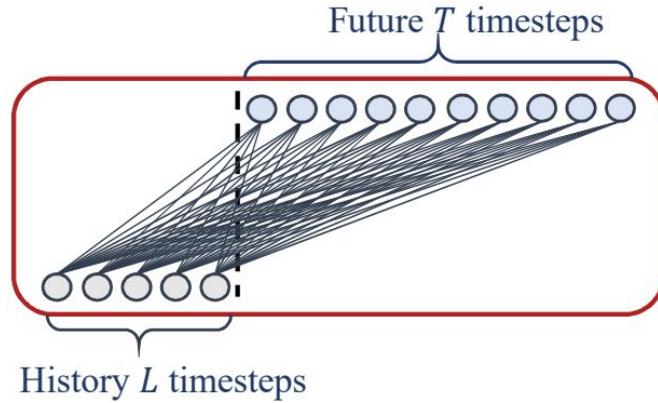
- no activation function
- in paper called just **Linear** (or Vanilla Linear)
- naturally performs direct multi-step (DMS) forecasts
- univariate, but performs surprisingly well for multivariate problems



Linear

"Are Transformers Effective for Time Series Forecasting?" A. Zeng et al.

- relies **exclusively** on order and magnitude of time series values
- similar to AR(p) model, but:
 - L is much larger than p , e.g. $L=96$
 - direct multi-step forecasts (DMS) reduce error
- models "how long ago something happened and how strongly"
- simple, which helps avoid overfitting and is very fast to train
- **just one** hyperparameter L , cheap & easy to tune



DLinear and NLinear

- expansions of the Linear model, from the same paper
- **DLinear (Detrended Linear):**
 - first detrends time series with moving average
 - uses 2 Linear models, for trend and remainder, forecast is their sum
 - performs better for data with clear trend
- **NLinear (Normalized Linear):**
 - first normalizes by subtracting the last value from time series
 - predicts normalized series, adds back value to forecast
 - just a normalization that should stabilize training

Linear model equivalence

"An Analysis of Linear Time Series Forecasting Models" W. Toner, L. Darlow

- authors prove that:
 - DLinear, NLinear and Linear are **equivalent** to OLS linear regression
 - NLinear is just Linear + constraint (rows sum to 1)
 - all equivalent models have **closed formula** from OLS
- uses SVD for training, which is great: optimal, fast, stable
- incredibly simple, but wins in 72% of experiments, and performs great
- just a single hyperparameter - **lookback window L**
- even L2 regularization is not required

Linear networks - pros and cons

Pros:

- closed formula OLS
- great performance
- very fast, stable, simple
- avoids overfitting
- just a single hyperparameter

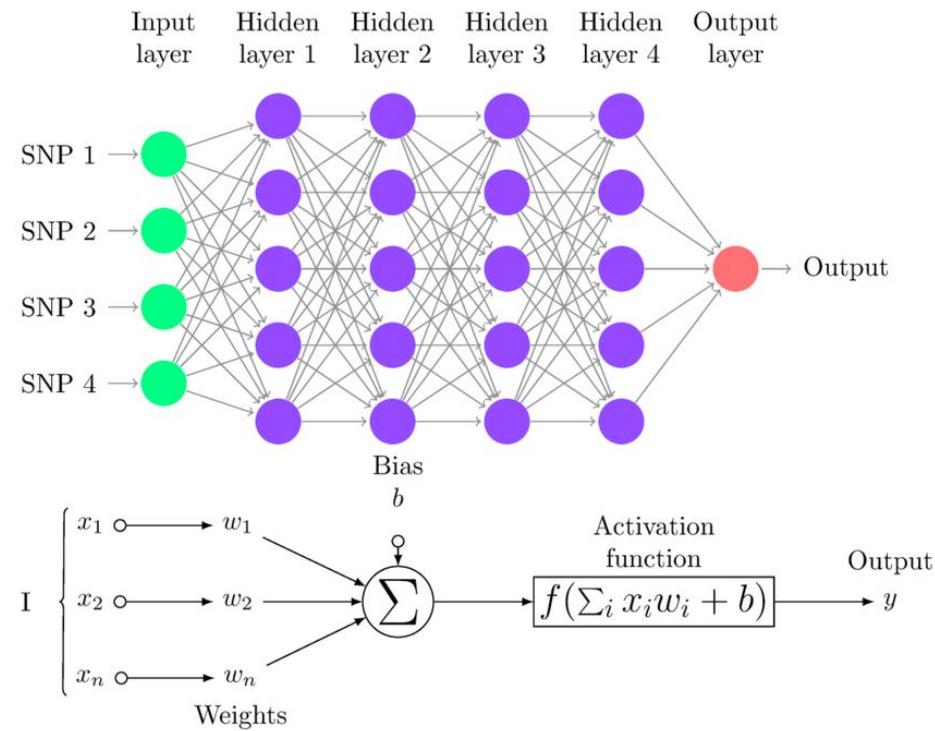
Cons:

- cannot learn very complex relations
- univariate - performs worse for strong cross-series correlations
- requires long time series for larger lookback L and learning long relations

MLP-based models

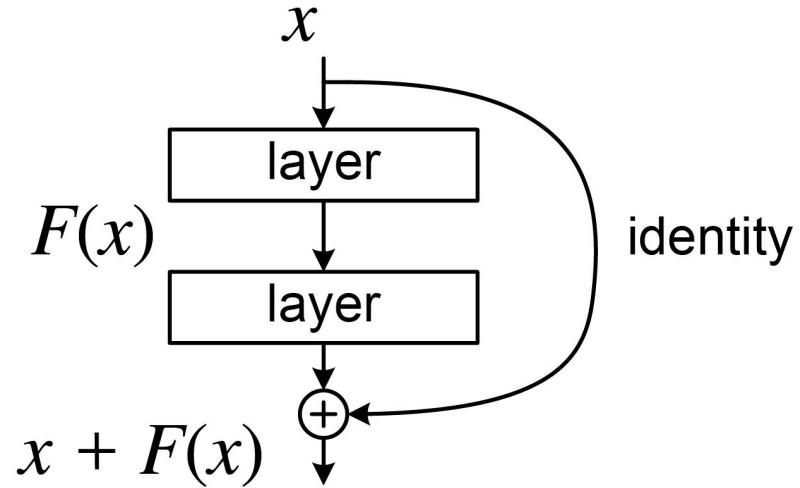
MLP - refresher

- multilayer perceptron
- built from **fully-connected layers**
$$X_{l+1} = \sigma(W^T X_l)$$
- non-linear **activation function** enables stacking layers, e.g.:
$$\text{ReLU}(x) = \max(0, x)$$
- can combine with anything differentiable



Residual networks - refresher

- **skip connections** in neural networks
- created as a simple engineering trick
- very commonly used
- advantages:
 - faster training
 - more stable loss
 - better results
- requires output to have same shape as input

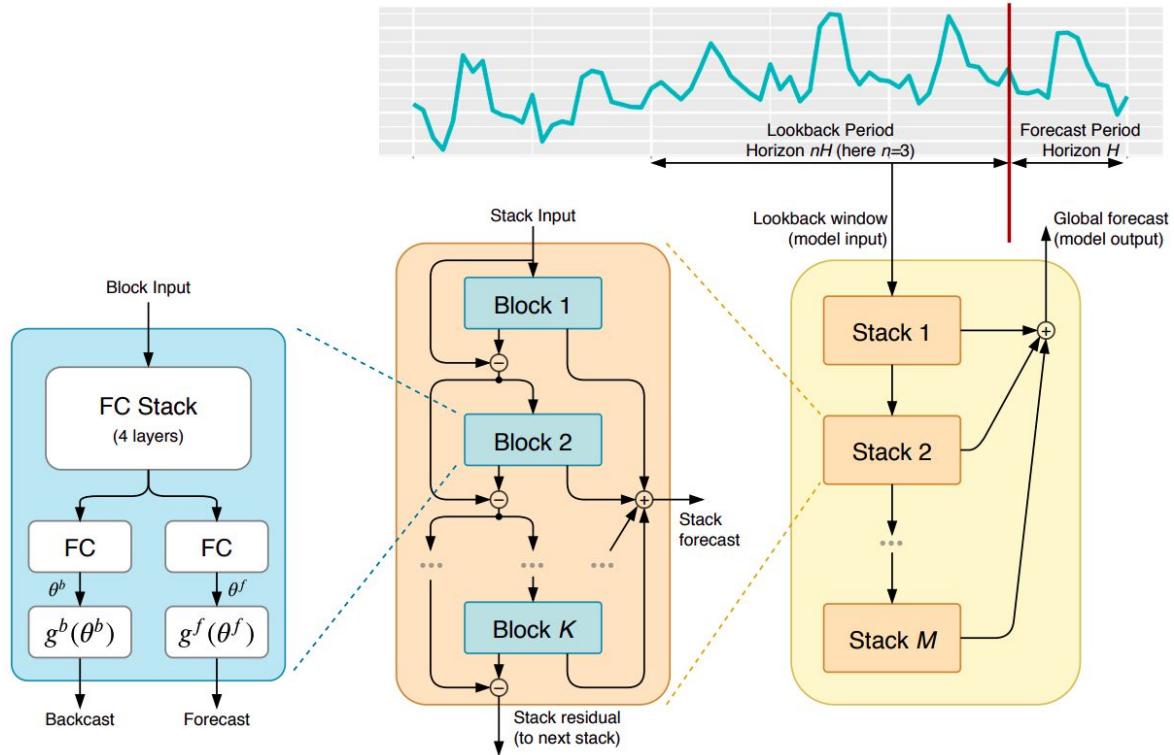


N-BEATS

"N-BEATS: Neural basis expansion analysis for interpretable time series forecasting"
B. Oreshkin et al.

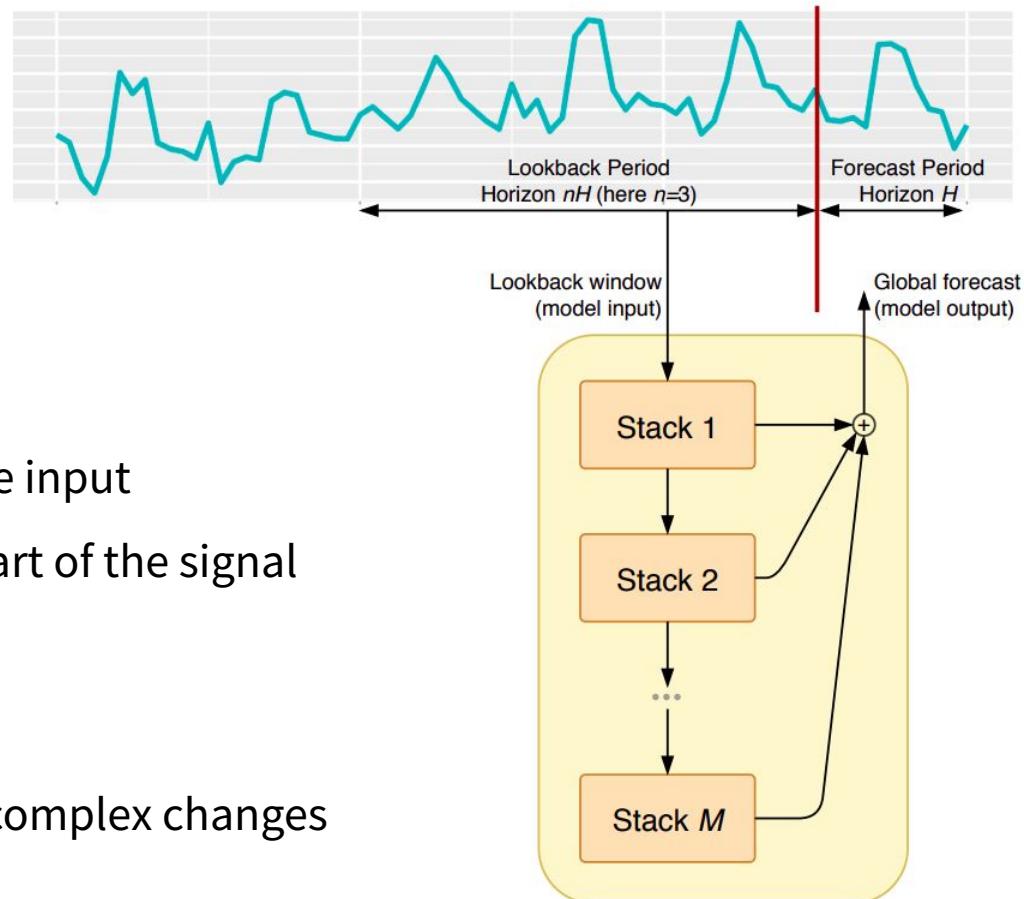
Combines a few key ideas:

- learn everything from raw data
- MLP as a basic building block
- doubly residual stacking
- stacked blocks
- basis expansion
- backcast



N-BEATS - stacked architecture

- first stack gets raw data of length nH
- model forecast is a sum from all stacks
- each stack has **2 outputs**:
 - forecast of length H
 - **residual** of its inputs
- stack subtracts what it learned from the input
- further ones only have to predict the part of the signal
- encourages **specialized** stacks, e.g.:
 - first learns trend - simple shapes
 - second learns seasonality - more complex changes

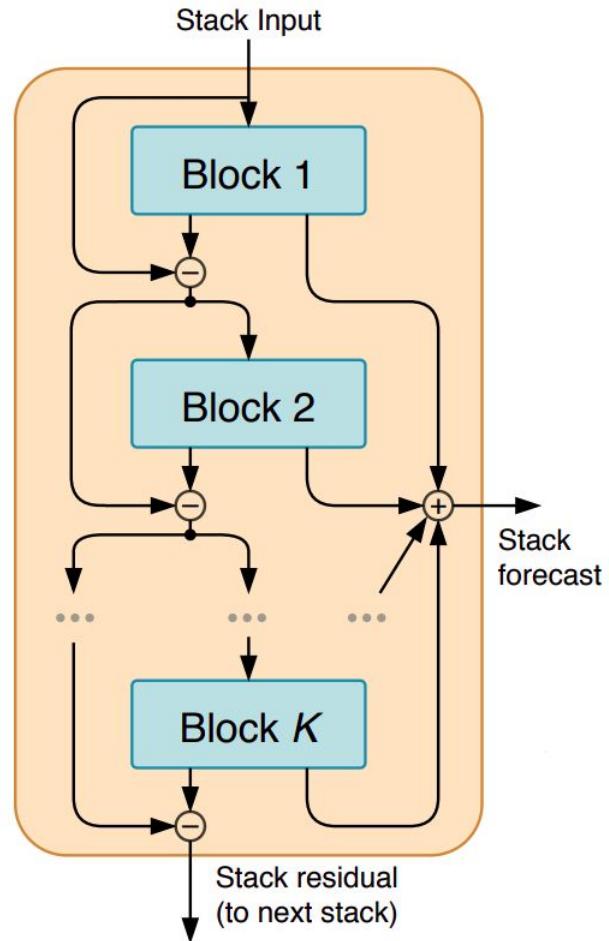


N-BEATS - variants

- paper proposed 2 variants: general and interpretable
- **general (N-BEATS-G):**
 - some N stacks, no structure enforced (linear basis)
 - free to learn arbitrarily complex relations, but need enough data
- **interpretable (N-BEATS-I):**
 - 2 stacks: trend and seasonality
 - trend stack: outputs trend forecast & detrended data
 - seasonality stack: outputs seasonality forecast
 - use dedicated basis functions for inductive bias:
 - polynomial (trend)
 - Fourier (seasonality)

N-BEATS - stack

- similar idea inside each stack, but with **blocks**
- block has 2 outputs:
 - partial forecast
 - **backcast**, estimating (reconstructing) its input data
- **residual connection:**
 - original data - backcast = residual
 - input into the next block
 - makes the job easier - removes parts of signal (data)
- called **doubly residual learning** in the paper (for stacks and for blocks)
- stack forecast = sum of block forecasts

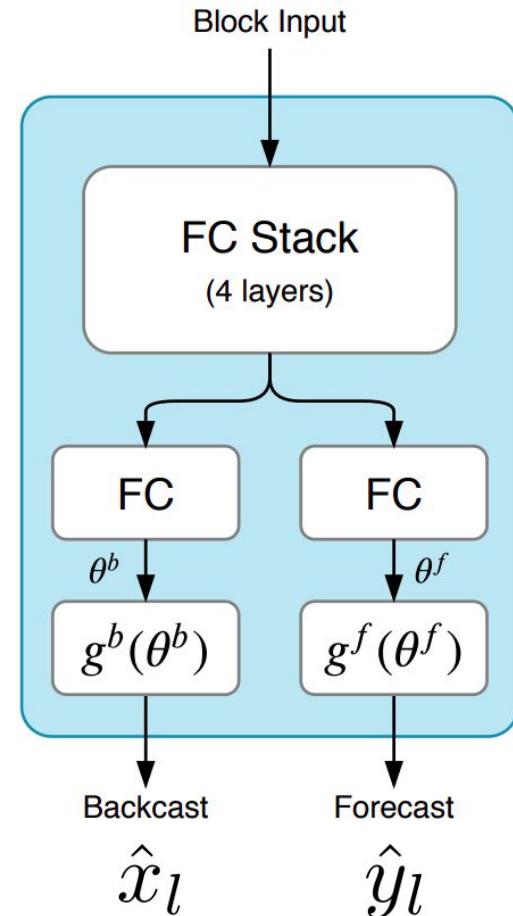


N-BEATS - block

- 2 outputs: forecast \hat{y}_l and backcast \hat{x}_l
- major idea is to predict **basis coefficients** of basis g
- allows encoding **inductive bias** through basis choice, e.g. seasonality is periodic
- forecast and backcast use the same basis, but separate weights
- generic architecture uses linear basis, which just matrix multiplication (linear projection)

$$\hat{y}_l = W_f \theta_f + b_f$$

$$\hat{x}_l = W_b \theta_b + b_b$$



N-BEATS - block

- interpretable variant encodes information in basis choice
- trend is uses **polynomial basis** of low degree, in paper $p=2$:

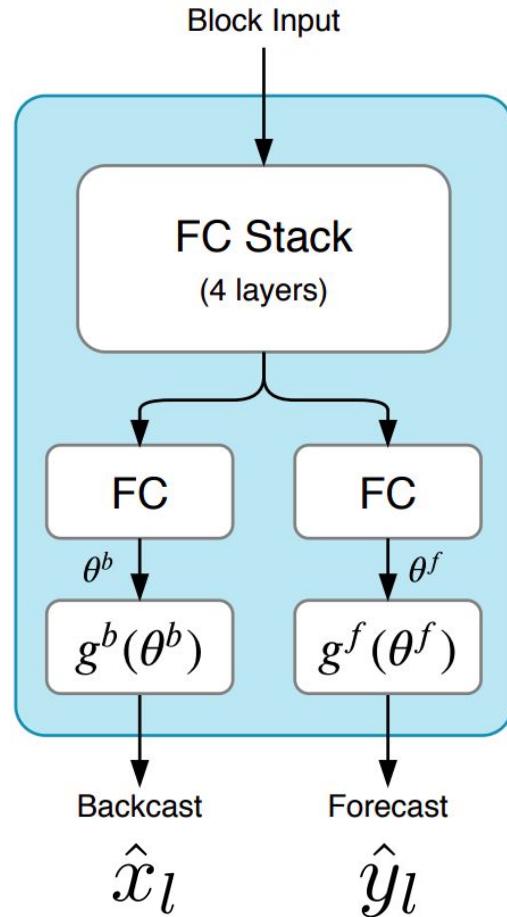
$$\hat{y}_l = \sum_{i=0}^p \theta_{f,i} t^i$$

t^i - time steps vector, linear grid raised to a given power

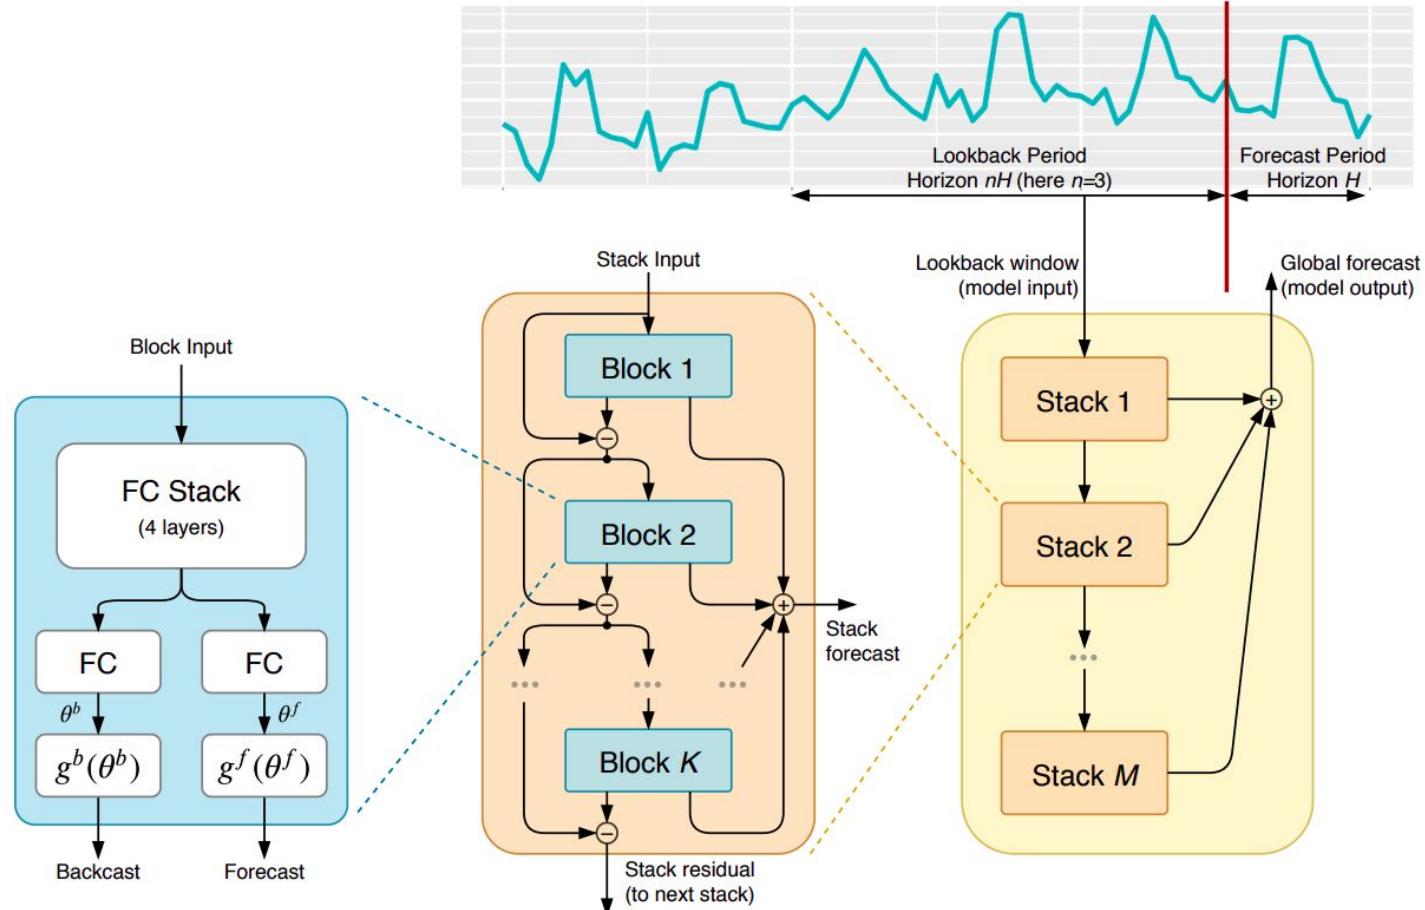
$$t^i = [0, 1, 2, \dots, H-2, H-1]^i / H$$

- seasonality uses **Fourier basis**

$$\hat{y}_l = \sum_{i=0}^{\lfloor H/2-1 \rfloor} \theta_{f,i} \cos(2\pi i t) + \theta_{f,i+\lfloor H/2 \rfloor} \sin(2\pi i t)$$



N-BEATS - recap



N-BEATS - pros and cons

Pros:

- very flexible
- good results
- can model very complex seasonality
- interpretable variant

Cons:

- only univariate
- does not scale well to long forecasting horizons (but: N-HiTS)
- no exogenous variables (but: N-BEATSx)

N-BEATS - additional resources

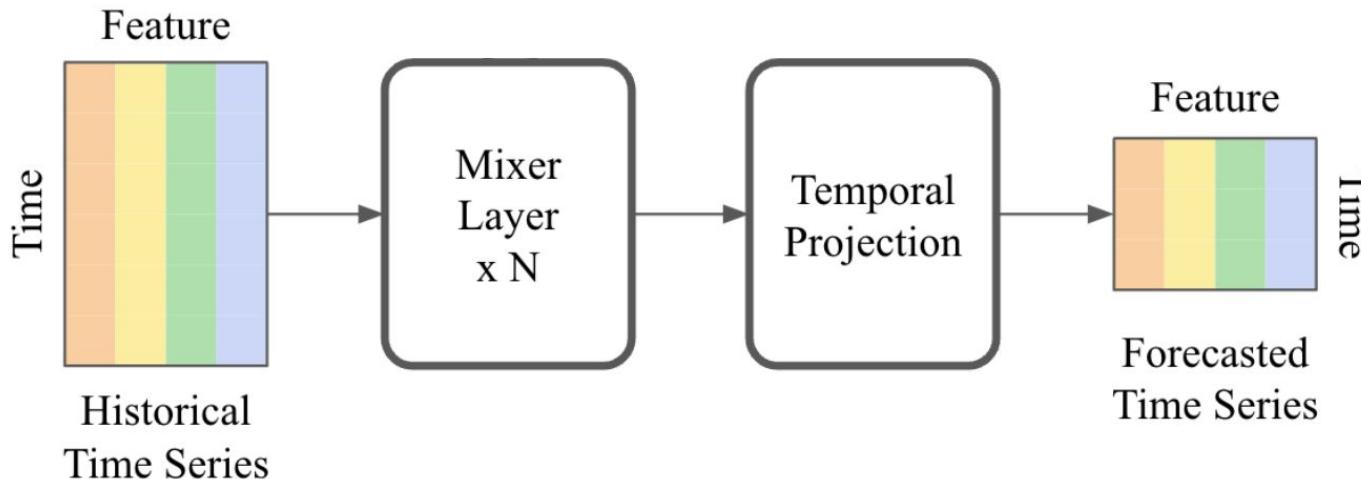
- "N-HiTS: Neural Hierarchical Interpolation for Time Series Forecasting" C. Challu et al.
- "Neural basis expansion analysis with exogenous variables: Forecasting electricity prices with NBEATSx" Kin Olivares et al.
- N-BEATS code in PyTorch Forecasting: blocks, whole model
- alternative explanation:
 - "N-BEATS — The First Interpretable Deep Learning Model That Worked for Time Series Forecasting" J. Dancker
 - "Optimizing Time Series Forecasting: Exploring N-BEATS Architecture for Improved Predictions" G. Sayago

TSMixer

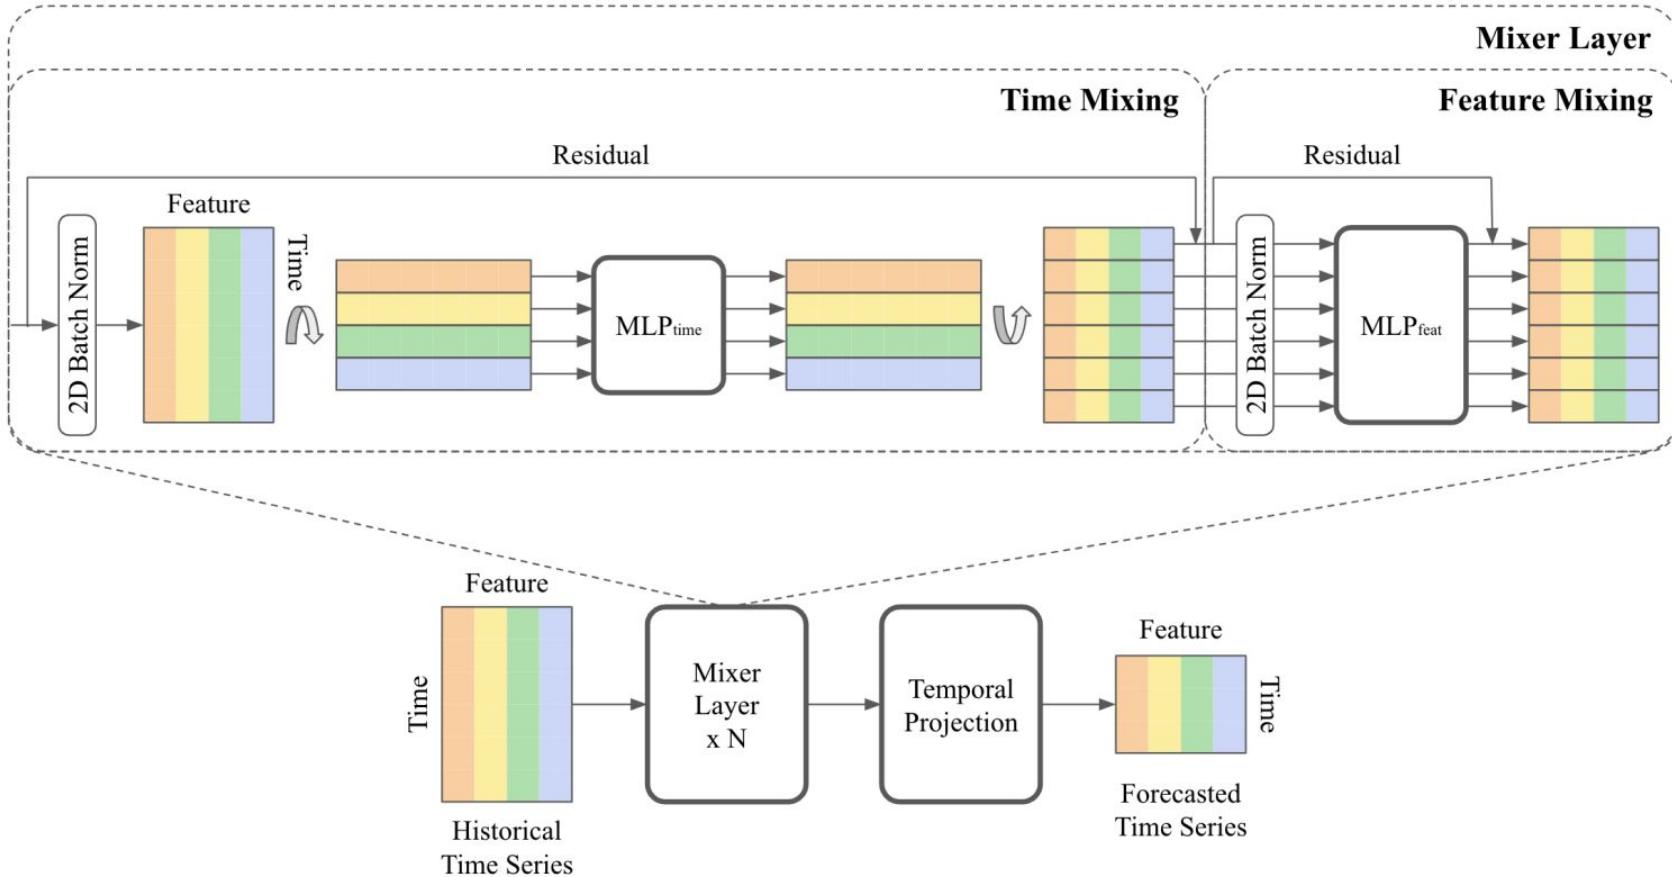
"TSMixer: An All-MLP Architecture for Time Series Forecasting" S. Chen et al.

- **idea:**

- "mixing" values in time or feature dimensions with MLPs
- process separately: univariate time, multivariate features, exogenous variables

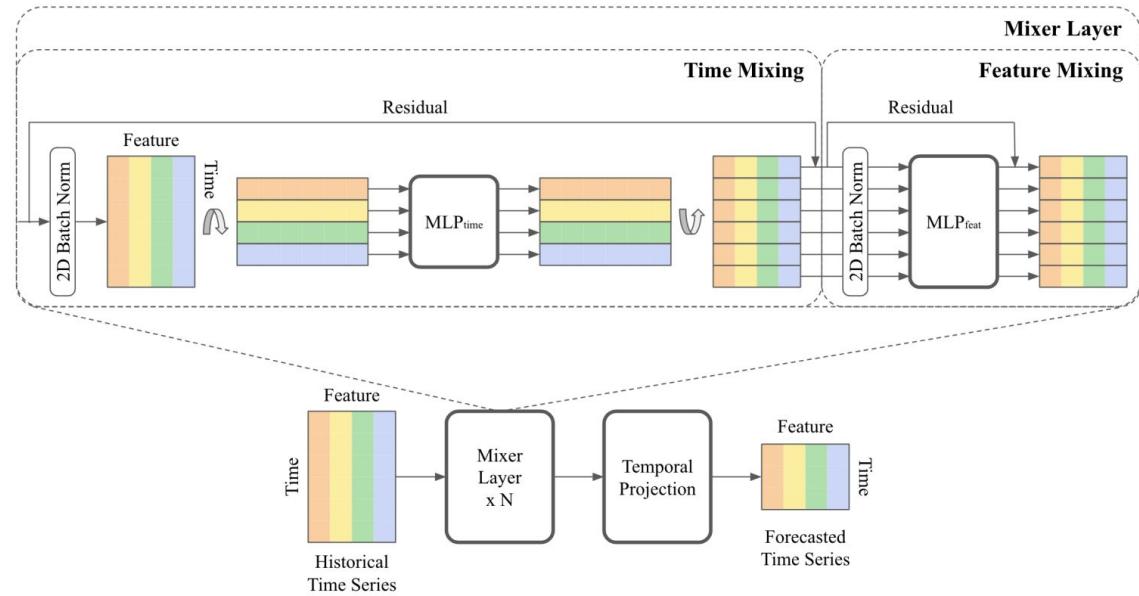


TSMixer - mixer layer



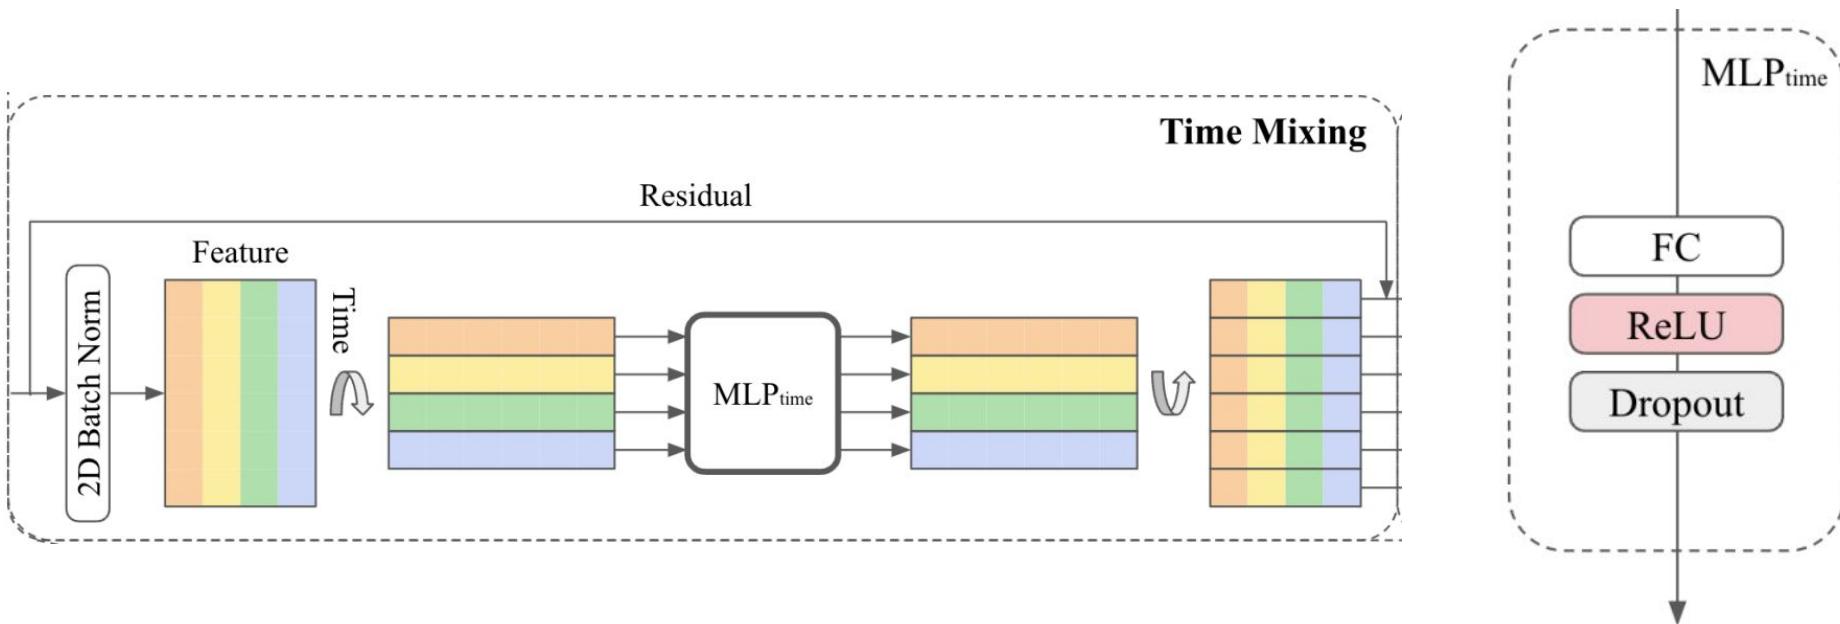
TSMixer - mixer layer

- **time mixing:** learn about time relations inside series
- **feature mixing:** learn about cross-series relations
- **separate** mixing reduces cost and complexity
- typical additions:
 - residuals
 - batch norm
- **temporal projection** is just a linear projection to horizon H



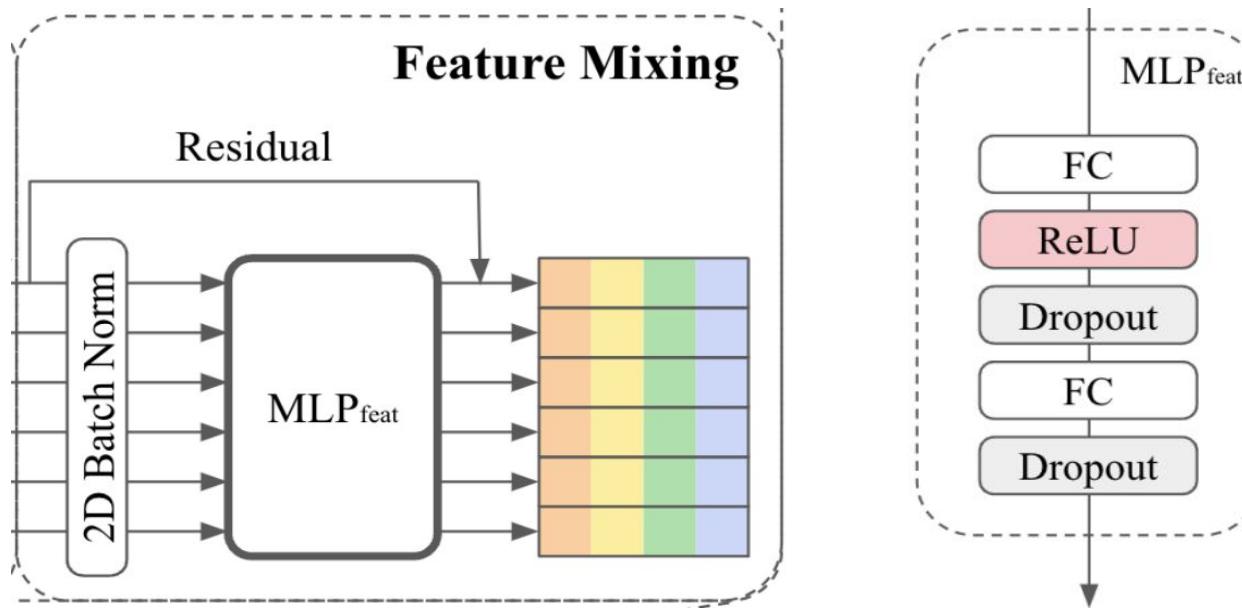
TSMixer - time mixing

- **inspired by** the Linear model
- the simplest non-linearity: 1-layer MLP
- extracts time-varying information inside a single time series



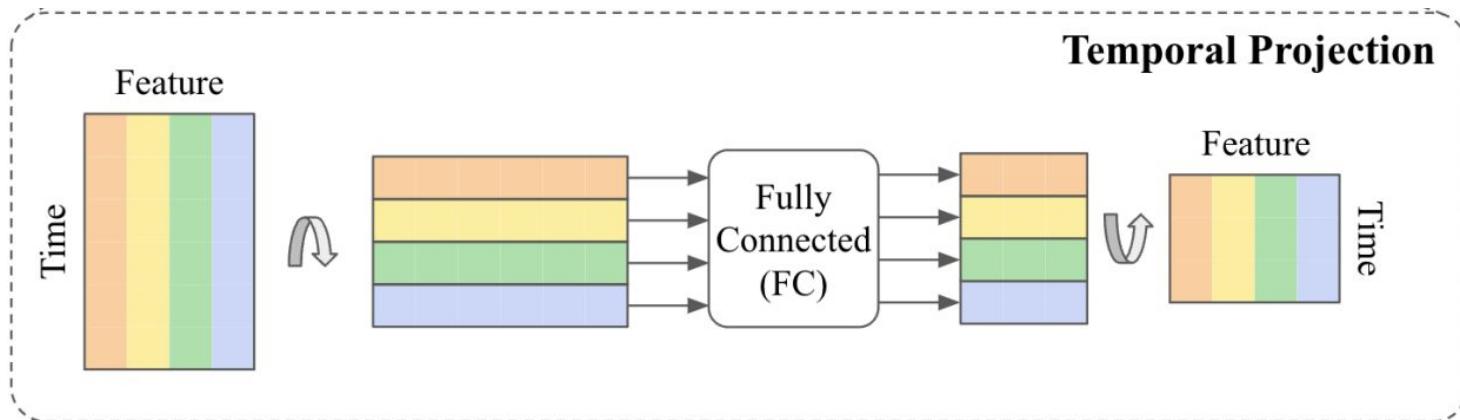
TSMixer - feature mixing

- **inspired by** the feature mixing in Transformer
- 2-layer MLP to learn more complex covariate relations
- extracts cross-series information



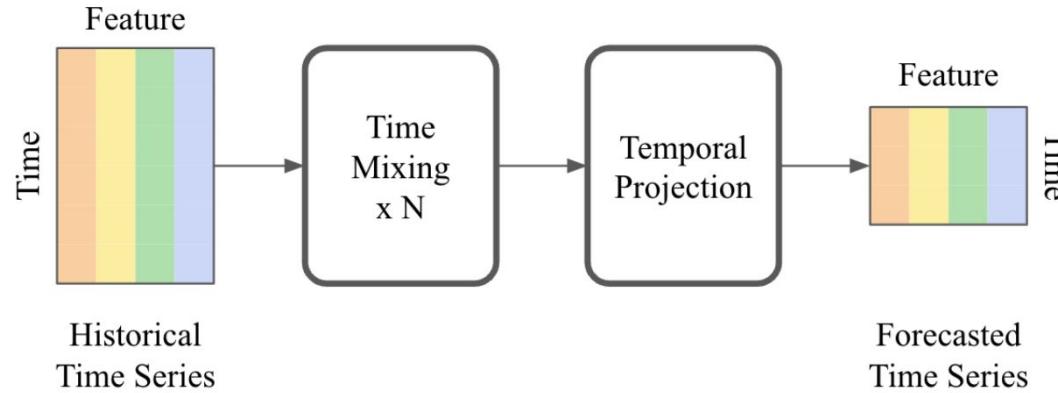
TSMixer - temporal projection

- **inspired by** the Linear model
- literally just a Linear model - simple linear projection from lookback L to to horizon H



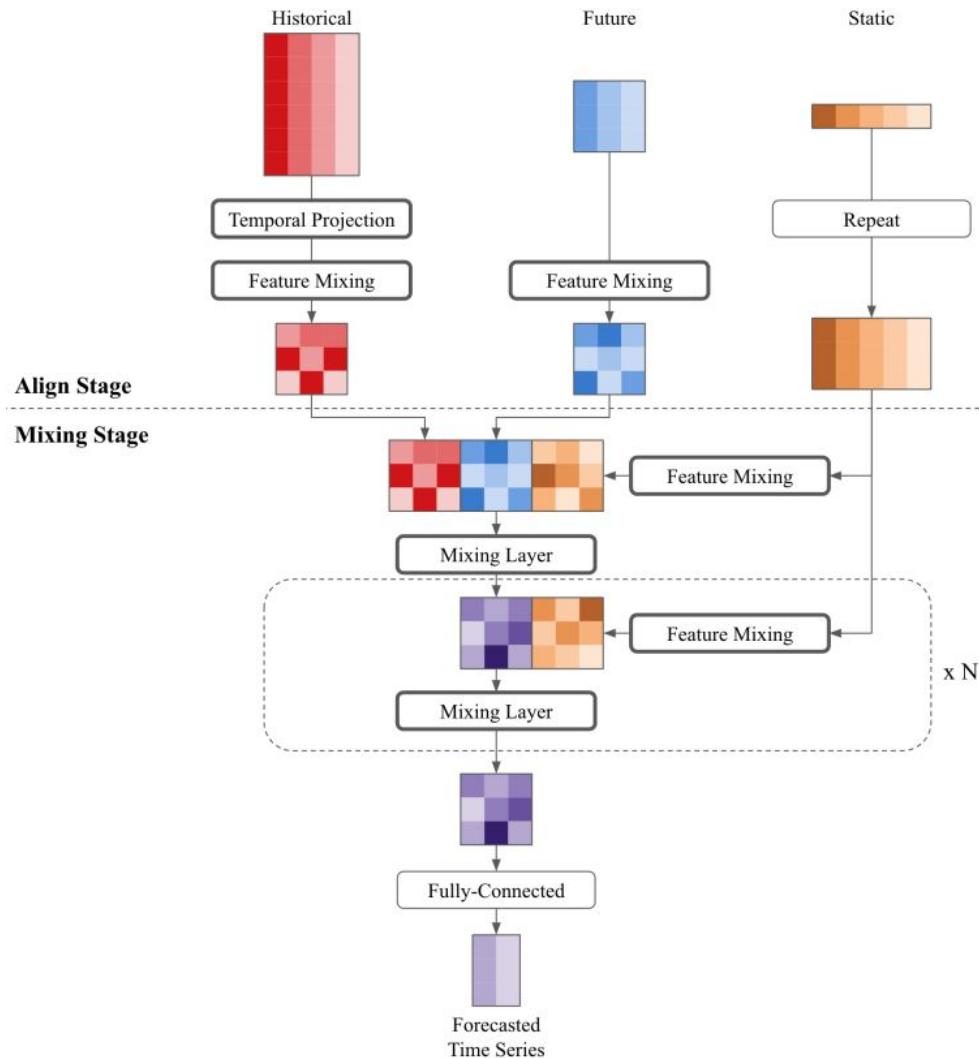
TMix-Only

- TSMixer variant for univariate time series
- only time mixing

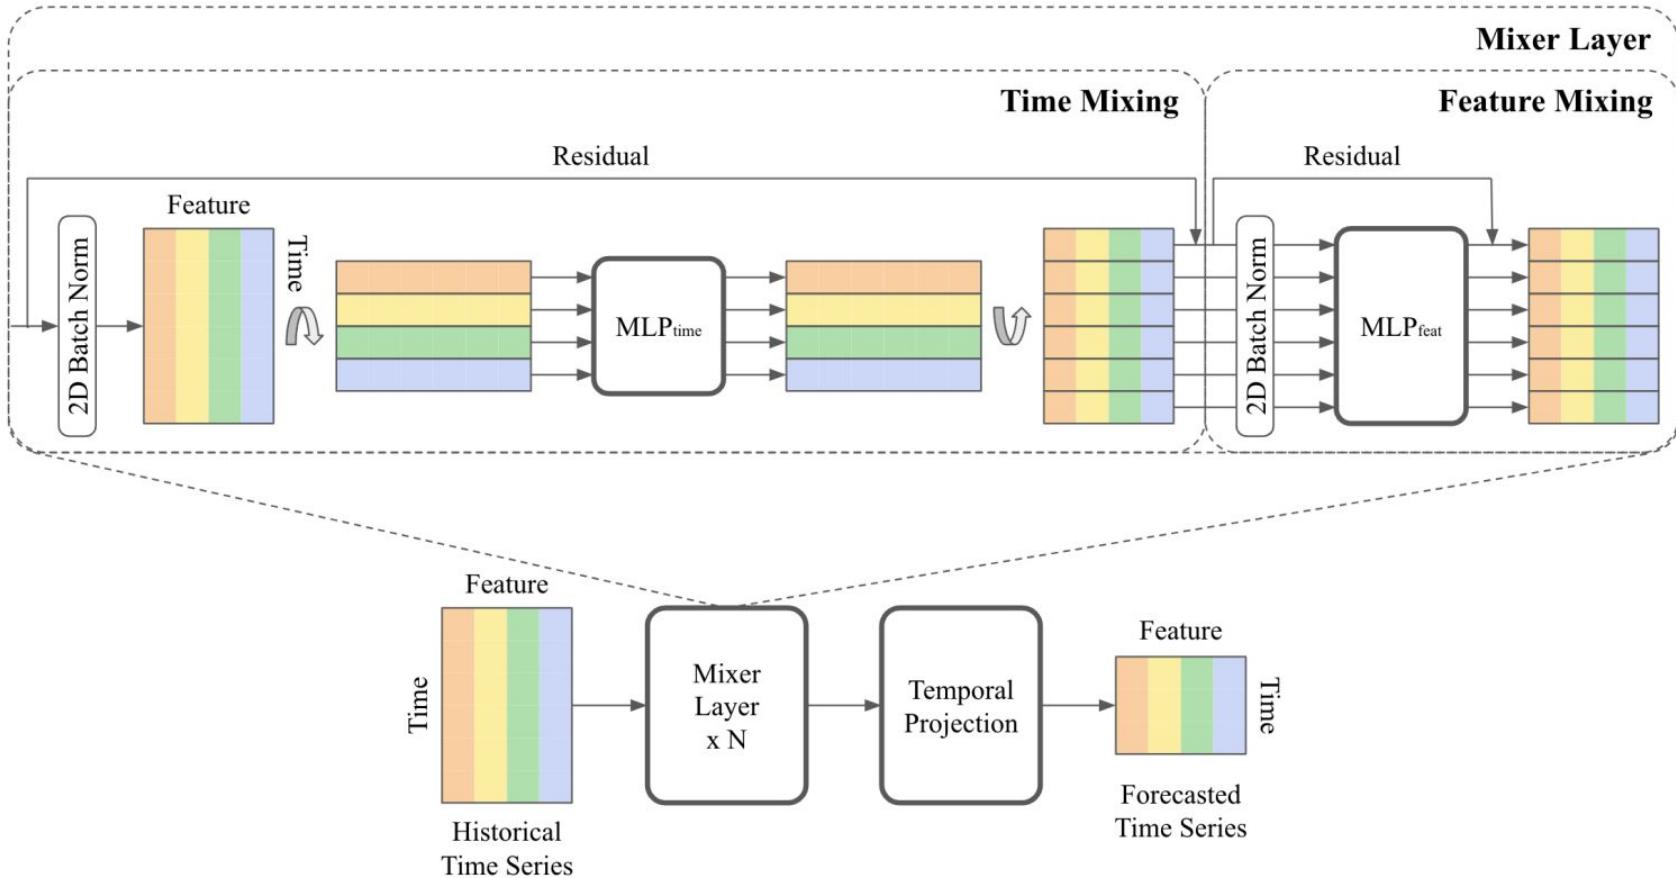


TSMixer - exogenous variables

- "auxiliary variables" in the paper
- can be:
 - **static**, e.g. shop location
 - **dynamic**, e.g. ongoing promotion
- historical data:
 - time series (1 or more)
 - dynamic auxiliary variables
- for more details, see the paper



TSMixer - recap



TSMixer - pros and cons

Pros:

- simple
- can model very complex relations
- uni- and multivariate
- exogenous variables support

Cons:

- can overfit with too little data
- computational cost (but not too high)
- not interpretable

TSMixer - warning!

- there are **two papers** with name "TSMixer"
- we talked about the one by Google!
- but there is also a one by IBM:

"TSMixer: Lightweight MLP-Mixer Model for Multivariate Time Series Forecasting" V. Ekambaram et al.

- Google one is better known, much simpler, people generally mean that one

Other interesting MLP-based models

- TimeMixer:

"TimeMixer: Decomposable Multiscale Mixing for Time Series Forecasting" S. Wang et al.

- time series downsampling & multiscale structure
- differentiable trend-seasonality decomposition (borrowed from Autoformer)
- many different mixings

- TiDE:

"Long-term Forecasting with TiDE: Time-series Dense Encoder" A. Das et al.

- MLP-based encoder-decoder
- flexible: univariate, multivariate, with exogenous variables
- quite small and very fast

Other mixing architectures

- time series:
 - [TimeMixer](#)
 - [Tiny Time Mixers \(TTMs\)](#)
 - [U-Mixer](#)
- computer vision:
 - ["MLP-Mixer: An all-MLP Architecture for Vision" I. Tolstikhin et al.](#)
 - ["Patches Are All You Need?" A. Trockman, J. Kolter - ConvMixer](#)
- graphs:
 - ["A Generalization of ViT/MLP-Mixer to Graphs" X. He et al.](#)
- NLP:
 - ["pNLP-Mixer: an Efficient all-MLP Architecture for Language" F. Fusco et al.](#)

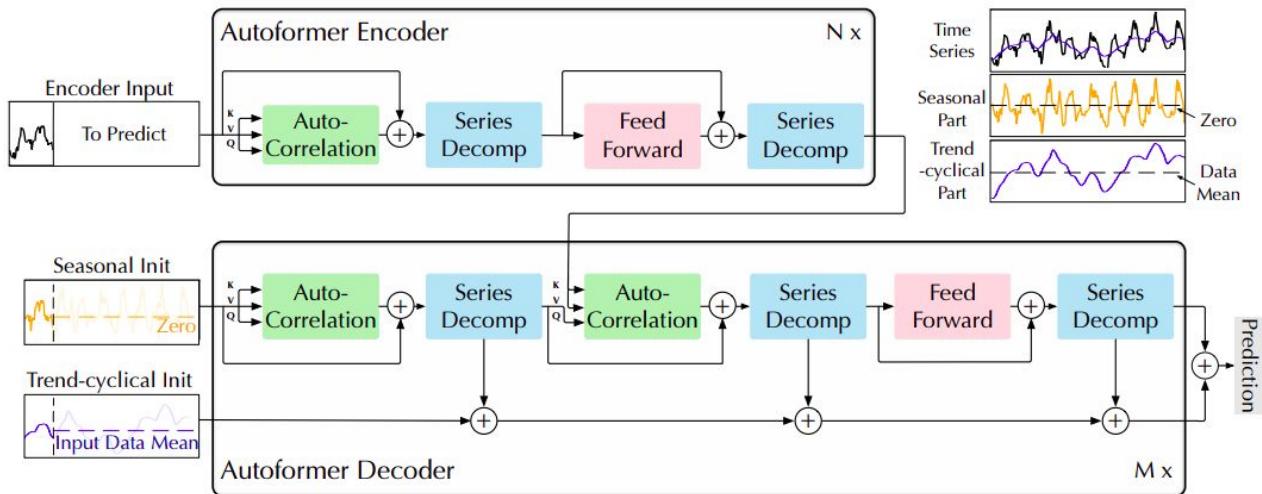
Transformers

Time series transformers

- **NLP-inspired transformer**, but with modifications for time series
- treat time series like sequence of words, task is to predict next words
- reduces cost and better learns time series information
- **quite varied:**
 - architecture: encoder-decoder / encoder-only
 - dimensionality: univariate / multivariate
 - pretraining: pretrained / trained from scratch (more frequent)
- foundation models are also based on transformers, but have visibly different trends in architecture - see further slides

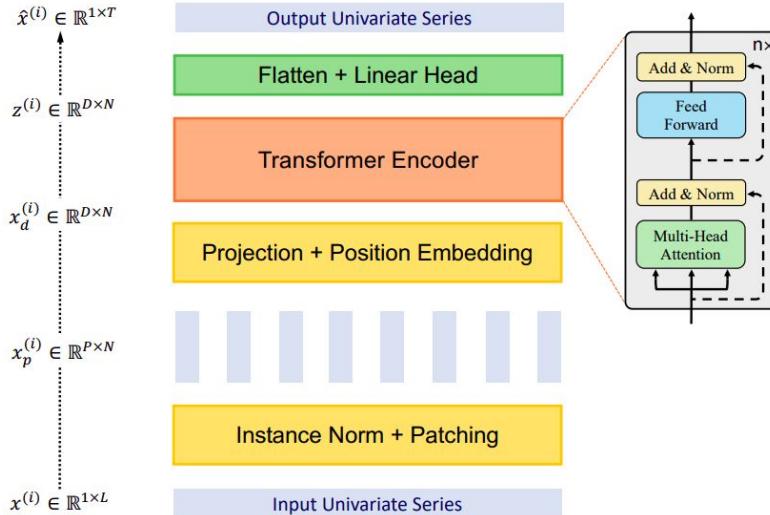
Encoder-decoder:

Autoformer



Encoder-only:

PatchTST



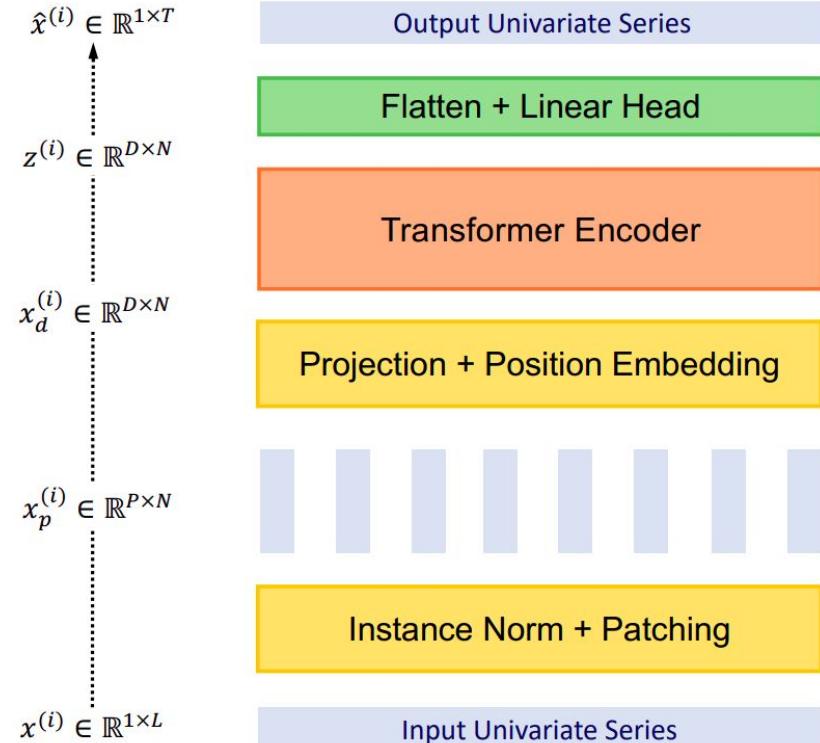
A bunch of transformers

| Model | Year | Architecture | Univariate / multivariate | Pretrained | Innovation |
|---------------|------|--------------|---------------------------|------------|---|
| LogTrans | 2019 | enc-dec | uni | no | LogSparse attention |
| Autoformer | 2021 | enc-dec | uni | no | Learn trend-seasonality and autocorrelation |
| Informer | 2021 | enc-dec | multi | no | Conv. subsampling, ProbSparse attention |
| FEDformer | 2022 | enc-dec | multi | no | Frequency enhanced attention |
| Pyraformer | 2022 | 2 variants | uni | 2 variants | Pyramidal attention with subsampling |
| NonStationary | 2022 | enc-dec | uni | no | Stationarization, De-Stationary attention |
| Crossformer | 2023 | enc-dec | multi | no | Hierarchical learning |
| PatchTST | 2023 | encoder-only | uni | yes | Patching, weight sharing, pretraining |
| iTransformer | 2024 | encoder-only | multi | no | Inverted tokenization (series-level) |

PatchTST

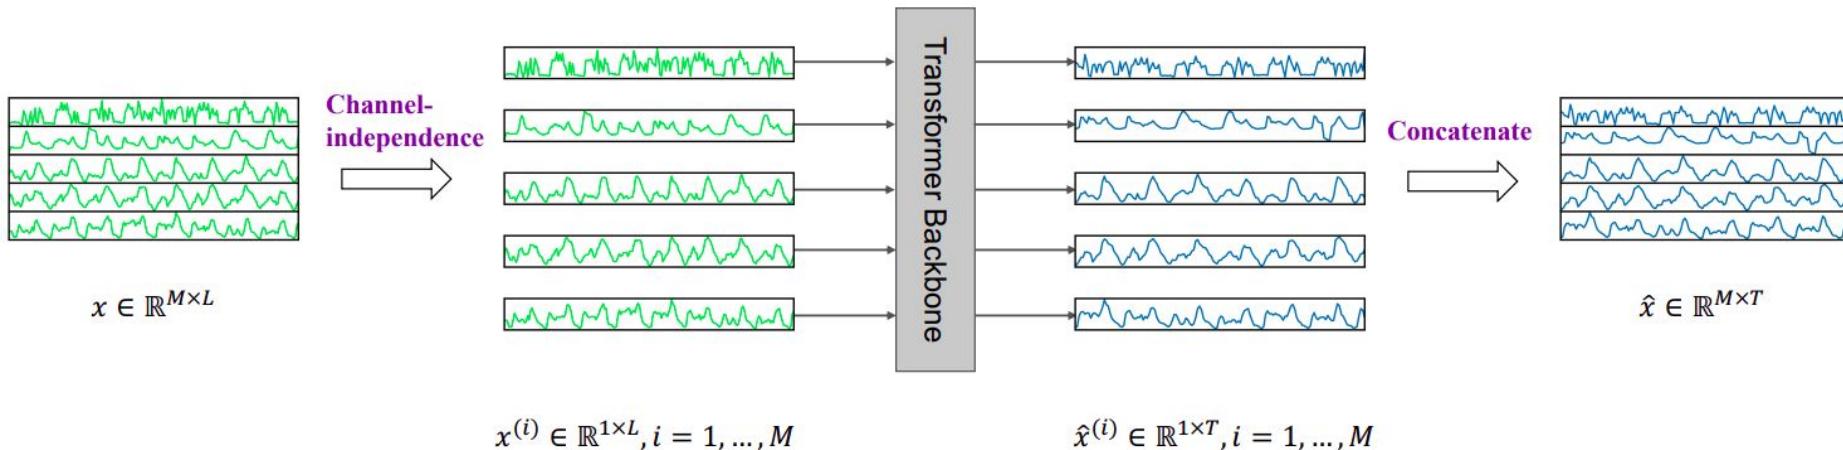
"A Time Series is Worth 64 Words: Long-term Forecasting with Transformers" Y. Nie et al.

- **encoder-only** transformer, very similar to BERT
- main idea is **patching** - tokenizing time series as "patches" of values
- after this, token sequence is just like a text sentence
- can pretrain with **masked modeling**, simply by masking and reconstructing patches
- uses previous 64 patches, which are equivalent to lookback $L=512$



PatchTST - channel independence

- univariate, processes time series separately - called **channel independence** in the paper
- but with **weight sharing** - all series have the same transformer!
 - regularizes, encourages better generalization
 - isolates effects from noisy channels
 - allows flexible number of time series

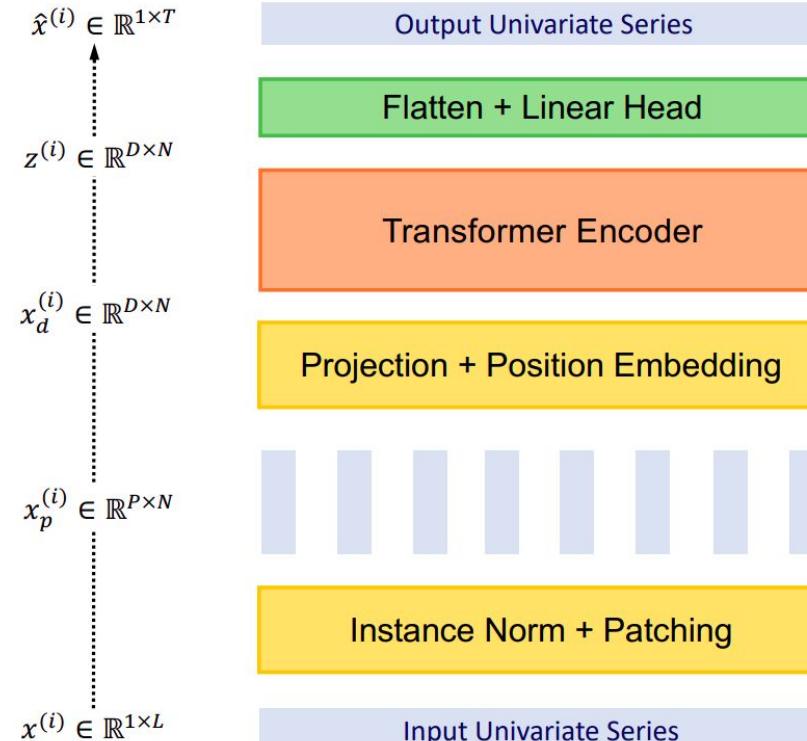


PatchTST - instance norm

- **instance normalization**
- time series standardization, before patching

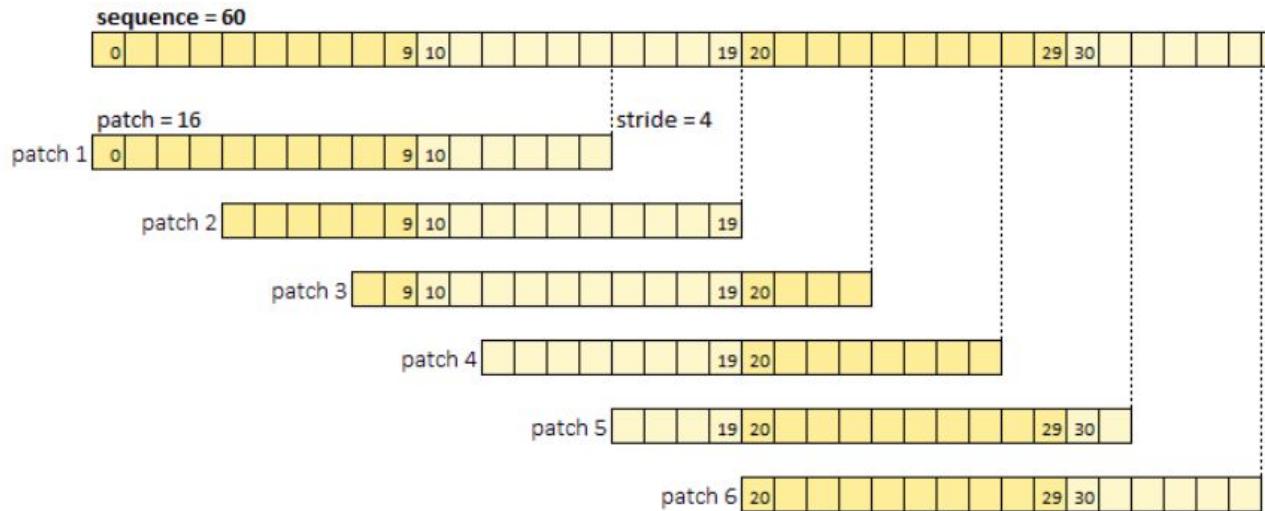
$$x' = \frac{x - \mu}{\sigma}$$

- processes time series separately
- reduces train/test **distribution shift**, making values distributions more similar
- added back to the final forecast



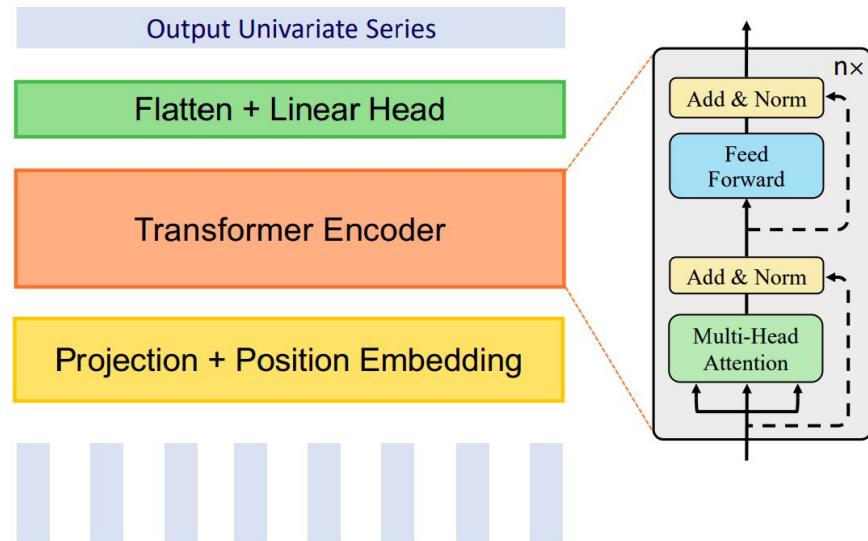
PatchTST - patching

- **patch** is a continuous series of values, creating a **token** for transformer layers
- greatly reduces complexity ($L \rightarrow L/S$ tokens) and allows longer lookback
- original paper models: patch length $P=16$, stride $S=8$
- **overlap** is similar to CNNs, slightly reduces overfitting due to shared data



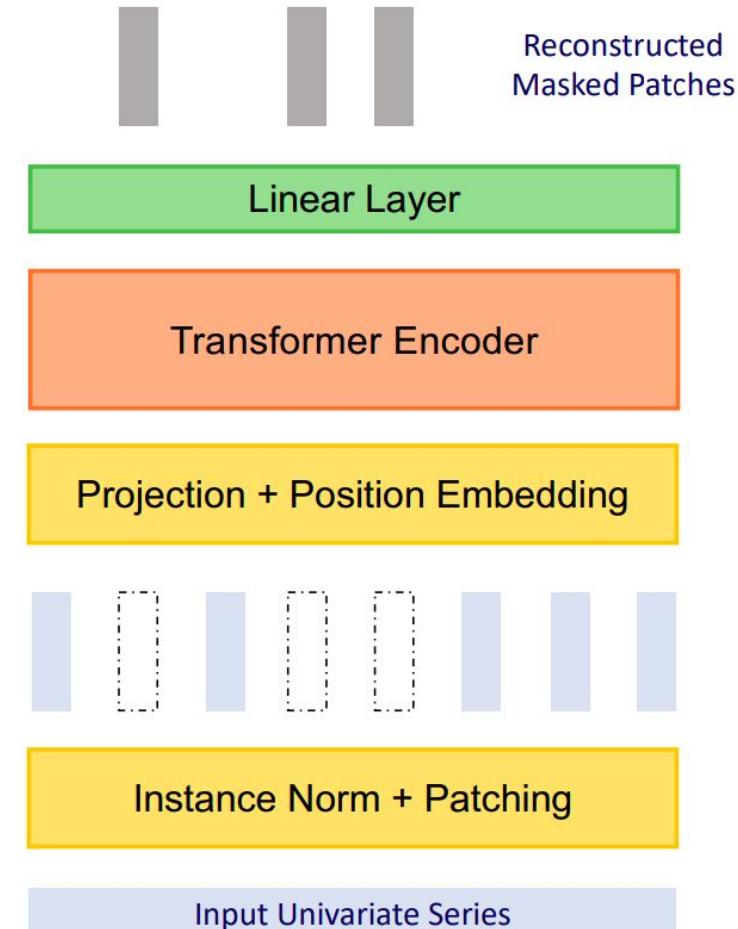
PatchTST - transformer

- similar to regular transformer layers, with a few changes
- **position embedding** encodes time and order
- multi-head self attention and 2-layer MLP use **GELU** activation
- **batch norm** instead of layer norm inside, it works better for time series
- paper parameters:
 - 3 layers
 - 16 attention heads
 - latent dimensionality D=128
 - MLP uses 128 → 256 → 128 dimensions

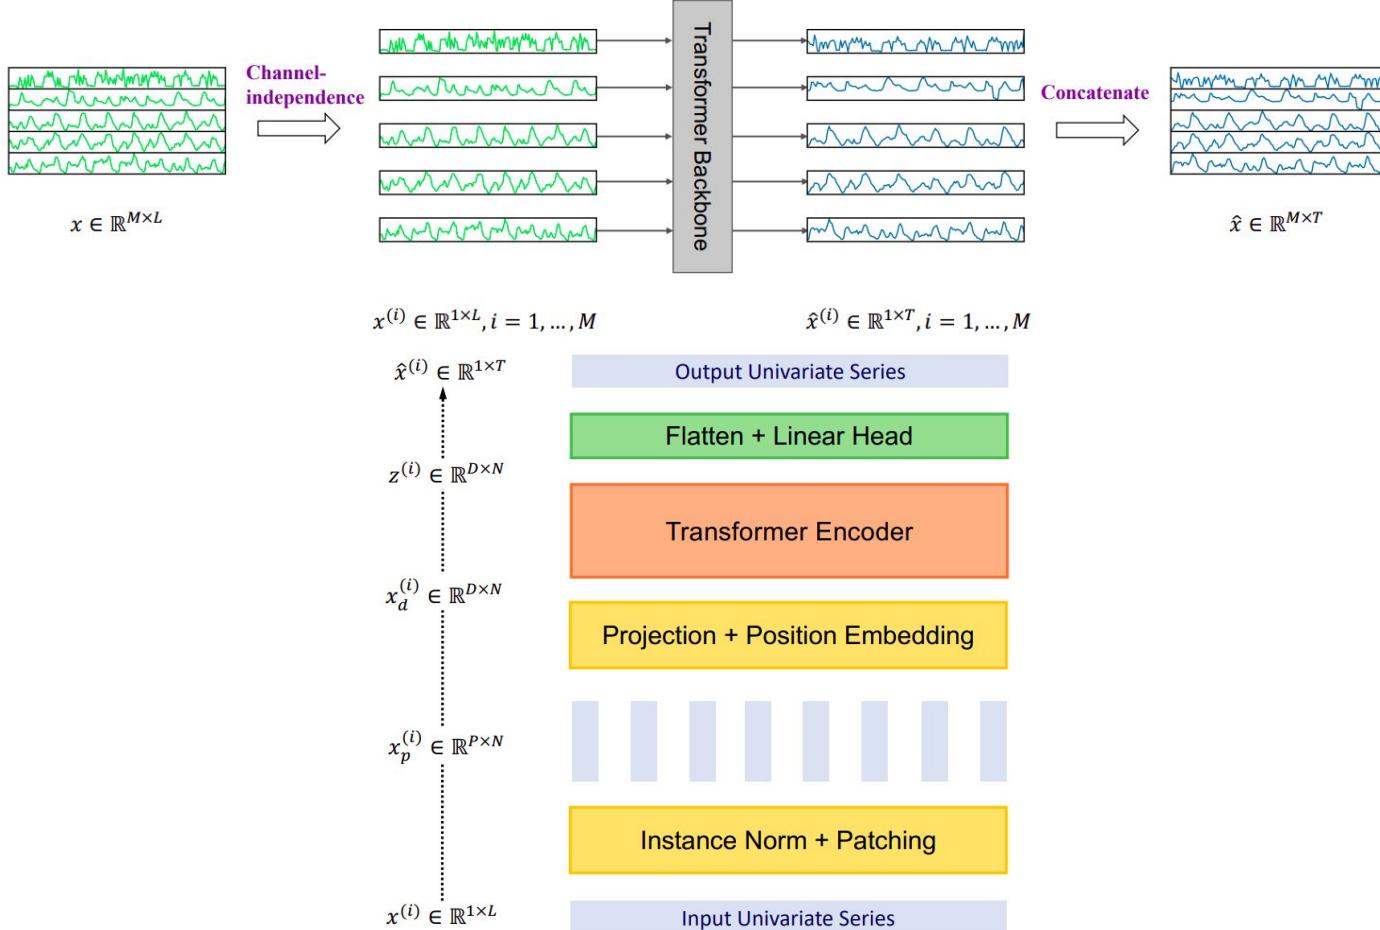


PatchTST - pretraining

- change head to **masked modeling**
- mask patches randomly and learn to reconstruct them, minimizing MSE
- non-overlapping patches to avoid data leakage, i.e. $P=S=16$
- typically results in better quality (not always!)
- but finetuning is always much faster than training from scratch
- can also finetune just the classification head, known as **linear probing**



PatchTST - recap



PatchTST - pros and cons

Pros:

- simple and fast (for a transformer)
- pretraining
- uni- and multivariate

Cons:

- easily overfits with too little data (but: linear probing)
- hard to pretrain effectively
- not interpretable
- no exogenous variables

Pretrained foundation models

Foundation models

- foundation model ([Wikipedia](#)):

"Model that is trained on broad data such that it can be applied across a wide range of use cases"

- **size matters:** large models + massive and diverse datasets + lots of computational power
- based on **representation learning:**
 - creating neural networks encoding general-purpose knowledge
 - internally create useful input representation (at least we hope so)
 - pretrain "domain expert", which can perform well on new tasks
- unique capabilities:
 - **few-shot learning** - with extremely short finetuning
 - **zero-shot forecasting** - no additional training, just input new data and get output

Time series foundation models

- **common features:**
 - transformers
 - quite simple architectures
 - patching (tokenization)
 - pretraining on massive datasets
- **varied:**
 - architecture: encoder-only, enc-dec, decoder-only
 - uni- / multivariate
 - exogenous variables support
- **idea:** rely on data and simple learning, rather than complicated models and handcrafted modules
- novelty: decoder-only, generative pretraining
- we will see if they are worthwhile in the future, for now those are a **research direction**

Time series foundation models - caution

- **be very cautious** when checking those models
- often made by companies, to create "hype" around "ChatGPT for time series"
- frequently not fully open source (incl. data, code, model weights)
- massive data requirement has its own challenges:
"Models work great when all test datasets are in your proprietary training dataset"
- whitepapers, preprints, technical notes etc. are **not peer-reviewed papers**

A bunch of time series foundation models

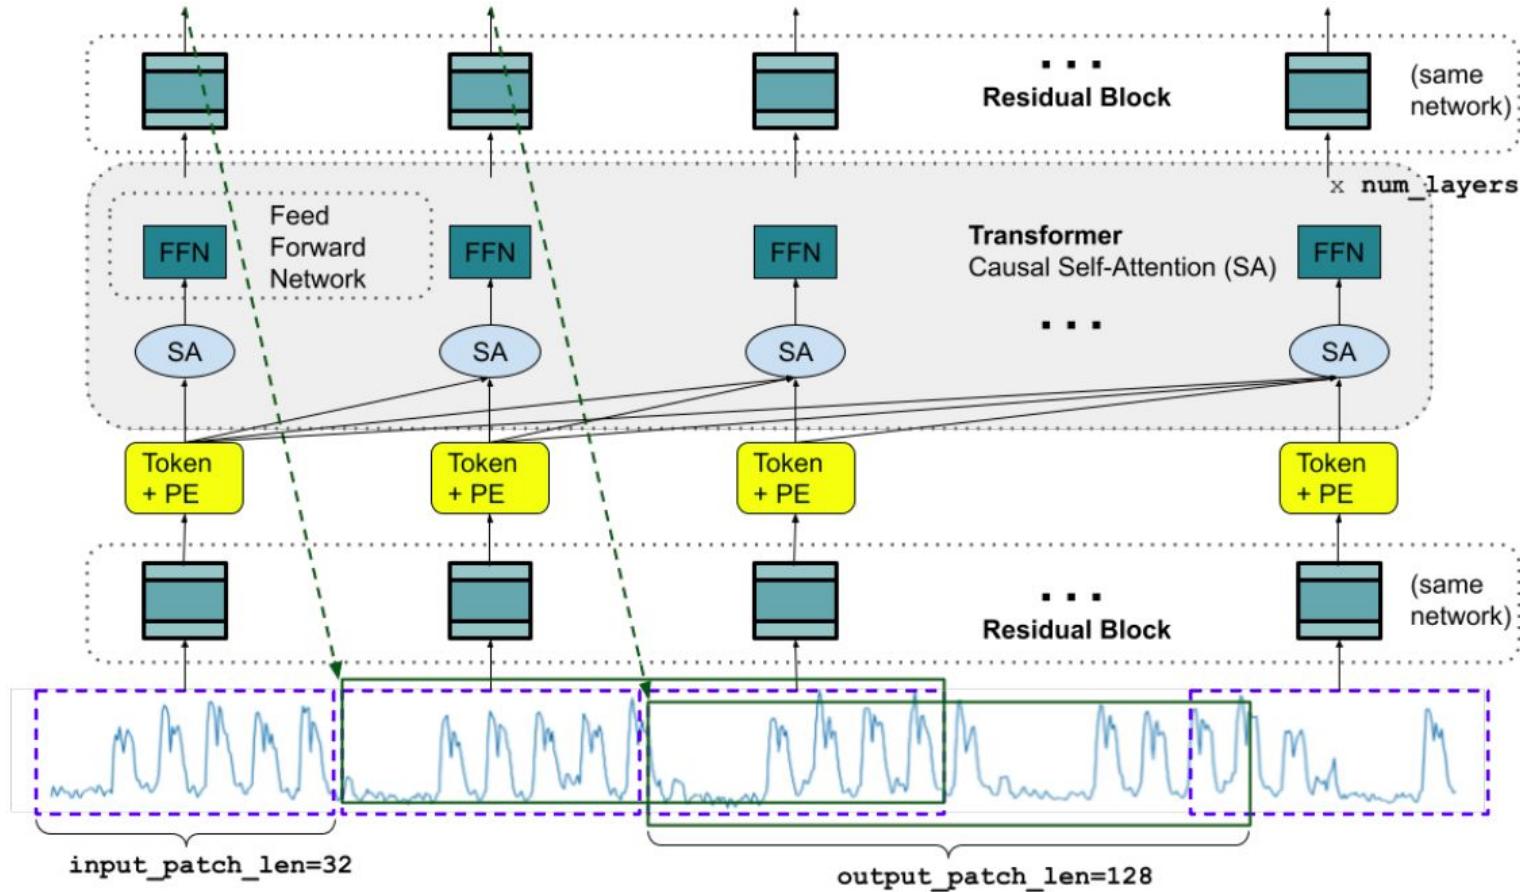
| Model | Year | Company / university | Open source? | Published paper? | Architecture | Univariate / multivariate |
|-----------|------|----------------------|--------------|---------------------------|--------------|---------------------------|
| TimeGPT | 2023 | Nixtla | no | no | enc-dec | uni |
| Lag-Llama | 2023 | Various (both) | yes | yes (NeurIPS workshop) | decoder | uni |
| TimesFM | 2024 | Google | yes | yes (ICML) | decoder | uni |
| Chronos | 2024 | Amazon | yes | no (TMLR reviews) | both | uni |
| Moirai | 2024 | Salesforce | yes | yes (ICML) | encoder | both |
| UniTS | 2024 | Harvard & MIT | yes | no | encoder | multi |

TimesFM

"A decoder-only foundation model for time-series forecasting" A. Das et al.

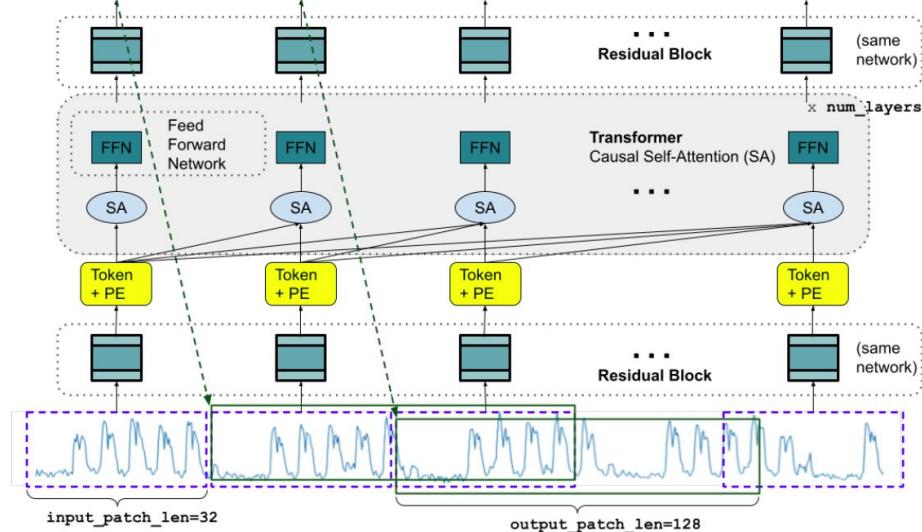
- Time series Foundation Model (TimesFM)
- **decoder-only** transformer, very similar to GPT
- **combination** of a few pretty simple ideas:
 - patching (tokenization)
 - decoder-only, generative pretraining
 - reasonable masking strategy
- in addition, they created a **massive pretraining dataset**, combining, e.g. Google Trends, Wikipedia page views, M4 datasets
- also used **synthetic**, generated datasets, exposing the model to different trends, shocks, seasonalities etc.

TimesFM - architecture



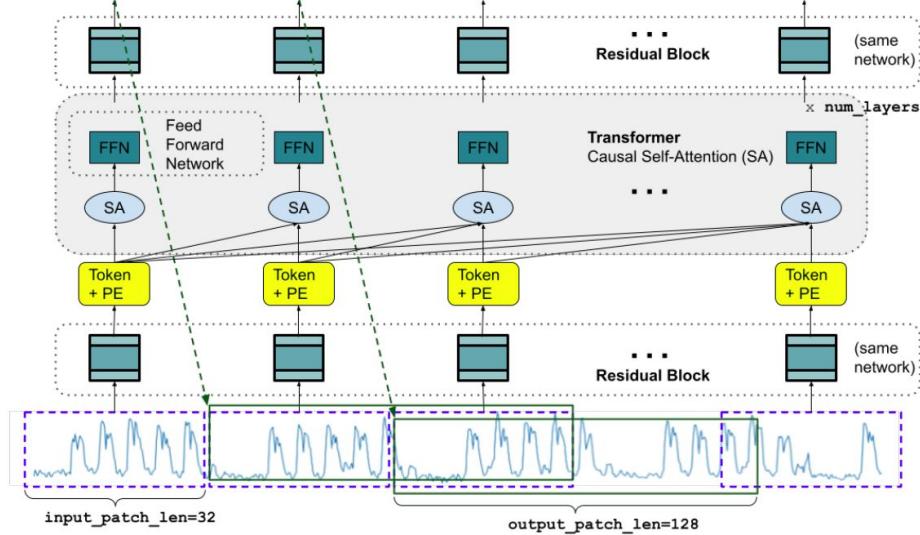
TimesFM - architecture

- basically, a standard **GPT transformer**
- **autoregressive** model that also makes direct multi-step (DMS) forecasts
- input patches are the "prompt"
- generates **long patches** at once, much longer than inputs
- this greatly reduces autoregressive error accumulation



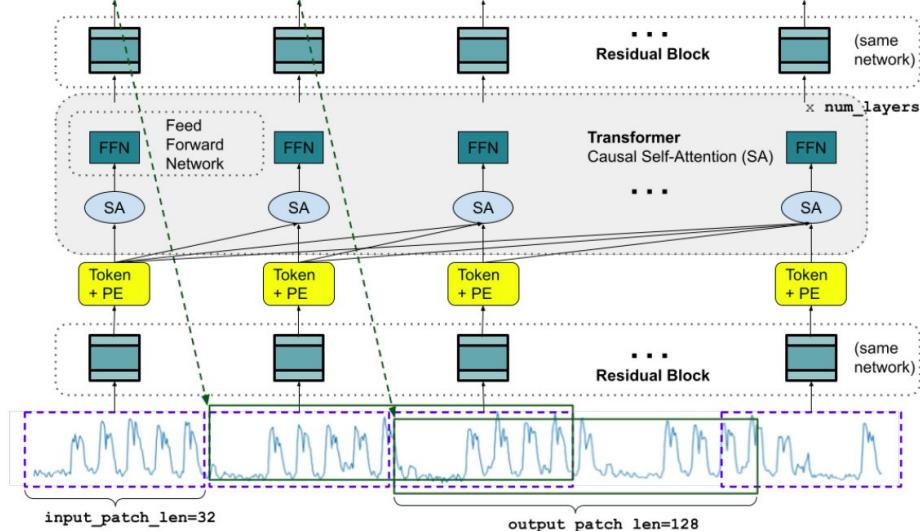
TimesFM - inputs and outputs

- "prompt" = **context**
- forecast = generated "words"
- sequence of **non-overlapping** patches
- input can have variable length, e.g. 32, 64, ..., padded when necessary
- **published model:**
 - input patch length 32
 - max context length 512 (16 tokens)
 - output patch length 128
- autoregressive forecast makes "full" steps, i.e. 128 here



TimesFM - transformer

- quite standard transformer
- causal attention with masking
- published model is quite **wide & deep**:
 - 20 layers
 - 16 attention heads
 - hidden size 1280 in all layers
 - 200M parameters in total
- generates 128-element vectors of floats



TimesFM - training

- just a regular training, with one detail - **patch masking**
- **problem:** for naive patches, model might learn to predict well only for context that is multiple of input patch length (e.g. 16, 32, 64, ...)
- **patch masking:**
 - for each time series in a batch, get a random number r from $[0, p-1]$
 - mask first r elements of the first patch, reducing the context
 - do this enough times and model will see all possibilities
- example:
 - $p=32, r=4$
 - first context is 28 (32-4), second is 60 (28+32), and so on

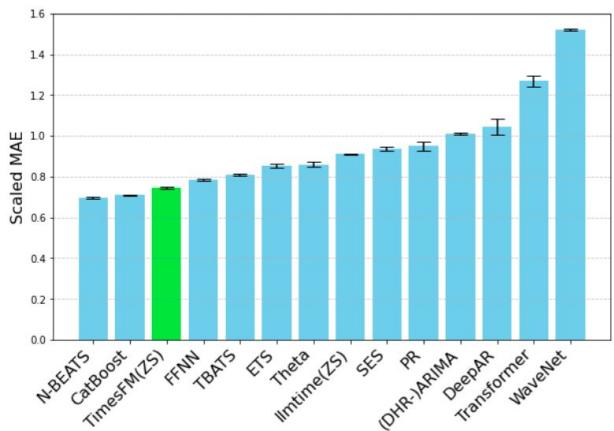
TimesFM - pretraining

- trained on a **mixture** of data
- a lot of **synthetic** data (20%):
 - trends, seasonalities
 - processes, e.g. ARMA
- real data (80%) chosen to give equal weights to different frequencies
- context depended on frequency:
 - 512 where possible
 - 256 for weekly
 - 64 for \geq monthly

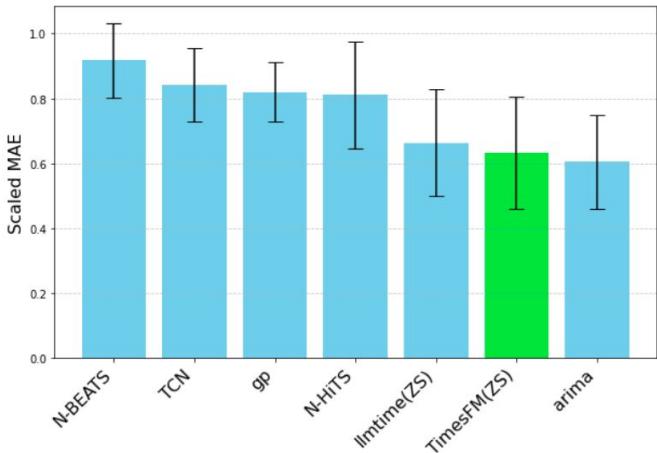
Table 1: Composition of TimesFM pretraining dataset.

| Dataset | Granularity | # Time series | # Time points |
|-------------------------------|-------------|---------------|-----------------|
| Synthetic | | 3,000,000 | 6,144,000,000 |
| Electricity | Hourly | 321 | 8,443,584 |
| Traffic | Hourly | 862 | 15,122,928 |
| Weather [ZZP ⁺ 21] | 10 Min | 42 | 2,213,232 |
| Favorita Sales | Daily | 111,840 | 139,179,538 |
| LibCity [WJJ ⁺ 23] | 15 Min | 6,159 | 34,253,622 |
| M4 hourly | Hourly | 414 | 353,500 |
| M4 daily | Daily | 4,227 | 9,964,658 |
| M4 monthly | Monthly | 48,000 | 10,382,411 |
| M4 quarterly | Quarterly | 24,000 | 2,214,108 |
| M4 yearly | Yearly | 22,739 | 840,644 |
| Wiki hourly | Hourly | 5,608,693 | 239,110,787,496 |
| Wiki daily | Daily | 68,448,204 | 115,143,501,240 |
| Wiki weekly | Weekly | 66,579,850 | 16,414,251,948 |
| Wiki monthly | Monthly | 63,151,306 | 3,789,760,907 |
| Trends hourly | Hourly | 22,435 | 393,043,680 |
| Trends daily | Daily | 22,435 | 122,921,365 |
| Trends weekly | Weekly | 22,435 | 16,585,438 |
| Trends monthly | Monthly | 22,435 | 3,821,760 |

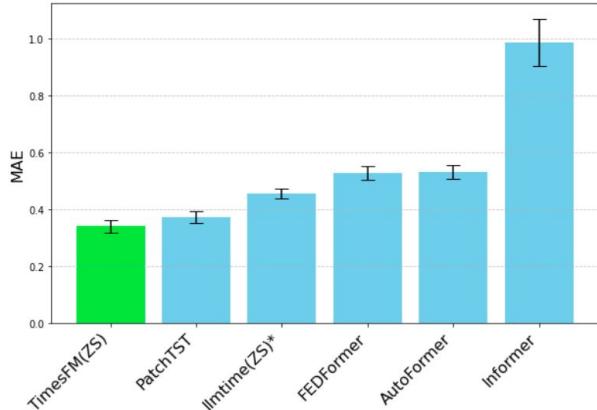
TimesFM - zero-shot results



(a) Monash Archive (Godahewa et al., 2021)



(b) Darts (Herzen et al., 2022)



(c) ETT (Horizons 96 and 192) (Zhou et al., 2021)

TimesFM - pros and cons

Pros:

- simple, yet powerful
- good results
- pretraining on a lot of data
- few-shot and zero-shot capabilities

Cons:

- computational cost
- only univariate
- no exogenous variables
- not interpretable

Questions?